

Deep Research Task 3A: Prior-Art and Evidence Mining for Known-Target Sign Validation

Executive Summary of the Evidence Landscape

The automated verification of American Sign Language (ASL) for targeted educational applications represents a distinct computational paradigm from generalized Continuous Sign Language Recognition (CSLR) or full Sign Language Translation (SLT). Because the educational application defines the target vocabulary in advance—providing a strictly bounded temporal window and a known semantic target—the machine learning objective shifts from open-vocabulary classification to known-target verification. This report exhaustively analyzes the prior art across Kaggle competition solutions, shipped commercial products, and the academic literature surrounding isolated sign language recognition (ISLR) and phonological assessment, extracting evidence on how to validate a learner's attempt at a known sign cheaply and in real time.

The evidence landscape reveals a profound bifurcation. In the domain of unconstrained, large-scale classification, highly optimized architectures utilizing 1D Convolutional Neural Networks (CNNs), Squeezeformers, and Spatio-Temporal Graph Convolutional Networks (ST-GCNs) have achieved remarkable inference efficiencies and robust accuracy ceilings on edge devices¹. Conversely, the explicit academic and commercial literature on "sign language pronunciation assessment" or "automated signing evaluation" is virtually nonexistent. Commercial applications frequently market camera-based feedback mechanisms, but empirical user reviews and technical audits expose these features as rudimentary video mirrors lacking active algorithmic validation⁴.

Consequently, the definitive template for known-target sign validation must be synthesized by retrofitting state-of-the-art spatial-temporal ISLR extractors with the scoring mechanics of speech Goodness-of-Pronunciation (GOP) algorithms. The findings indicate that lightweight 1D CNNs or ST-GCNs—when regularized heavily with temporal masking and Adversarial Weight Perturbation (AWP)—can reliably map continuous 3D skeletal data to dense phonological embeddings well within a 200-millisecond real-time latency budget. This report details the architectural parameters, augmentation recipes, and evaluation frameworks necessary to transfer these capabilities directly to a 100–300 sign curriculum verifier.

Section 1: Kaggle: Google – Isolated Sign Language Recognition (2023)

The Google Isolated Sign Language Recognition (ISLR) competition required predicting one of 250 signs utilizing MediaPipe Holistic landmarks, evaluated strictly under TensorFlow Lite on-device latency constraints. The winning solutions abandoned computationally expensive

vision transformers in favor of heavily regularized 1D CNNs and hybrid Conv-Transformer topologies².

Per-Source Record: Kaggle ISLR Top Solutions

Group / Author	Year	Task	Dataset	Input Modality	Metrics	Offline vs. Real-Time	Code/Weights Links
hoysong48 (1st Place)	2023	ISLR	Kaggle ASL Signs	MediaPipe Landmarks	LB 0.82	Real-Time (17ms TFLite)	²
Kolya Forrat (2nd Place)	2023	ISLR	Kaggle ASL Signs	80 Landmarks	LB ~0.80, CV 0.898	Real-Time Capable	⁶
Ruslan Grimov (3rd Place)	2023	ISLR	Kaggle ASL Signs	MediaPipe Landmarks	LB Top-tier	Real-Time Capable	⁷
ohkawa3 (4th Place)	2023	ISLR	Kaggle ASL Signs	MediaPipe Landmarks	LB Top-tier	Real-Time Capable	⁷
yuanzhe zhou (5th Place)	2023	ISLR	Kaggle ASL Signs	MediaPipe Landmarks	LB Top-tier	Real-Time Capable	⁷

Architectural Patterns of the Winners

The first-place solution utilized a 192-dimensional 1D CNN combined with a lightweight Transformer, totaling approximately 1.85 million parameters². The model ingested a dynamic temporal sequence up to a maximum length of 384 frames. The foundational logic for this architecture was the observation that inter-frame correlation in ASL is exceedingly strong; 1D CNNs extract these local temporal dependencies more efficiently than the global attention

mechanisms of pure Transformers².

The architecture began with a masking layer to handle variable-length inputs without relying on zero-padding during inference, ensuring strict train-test consistency². The 1D CNN blocks employed depthwise convolutions and causal padding to maintain precise mask indices throughout the network. The subsequent Transformer utilized Batch Normalization paired with the Swish activation function, eschewing the traditional LayerNorm and GELU combination. This specific architectural choice was driven by the necessity for a lighter inference footprint during TensorFlow Lite conversion, providing identical accuracy with reduced computational overhead².

The second-place solution approached the problem through an audio-spectrogram paradigm, utilizing an EfficientNet-B0 backbone⁶. The temporal landmark coordinates were interpolated to a fixed size of 160 frames and plotted against 80 preserved coordinate points, yielding a 160x80x3 tensor structurally analogous to a mel-spectrogram⁶. This demonstrates that 2D CNNs, typically reserved for image processing, can successfully model temporal pose sequences if the temporal dimension is treated as a spatial axis.

Landmark Selection and Preprocessing

Rather than ingesting the full MediaPipe whole-body graph, top solutions aggressively pruned the topological input space to reduce noise. The second-place solution retained exactly 80 points: 18 lip coordinates, 20 upper-body pose coordinates (arms, shoulders, eyebrows, nose), and all 42 hand coordinates⁶. A supplementary Transformer model in the same ensemble further reduced this to 61 points (40 lip points, 21 dominant hand points).

Normalization strategies varied but universally anchored the spatial data to a stable facial or torso node to achieve translation invariance. The first-place solution utilized the 17th MediaPipe landmark (the nose) as the absolute root, translating all coordinates relative to the nose center while maintaining dynamic inputs rather than dropping NaN values². Missing coordinates due to self-occlusion were intentionally filled with zeros post-normalization, effectively forcing the network to learn the presence of occlusion as a spatial feature rather than treating it as missing data⁶.

Augmentation and Regularization Recipes

Because the competition required training from scratch on a limited dataset, regularization was the most critical factor separating the top-tier solutions from the baseline. Training epochs frequently extended to 400 iterations, necessitating aggressive anti-overfitting measures².

The universal augmentation recipe deployed by top performers included spatial augmentations such as horizontal flipping (to simulate left-handed signers) and random affine transformations (scale, shift, rotate, shear)². Random cutout masking was applied heavily to spatial coordinates to simulate upstream pose extractor failures.

In terms of regularization, the first-place solution implemented Drop Path (stochastic depth) with a probability of 0.2 applied after each convolutional block, standard Dropout at a high rate of 0.8 prior to the global average pooling layer, and Adversarial Weight Perturbation

(AWP) with a lambda parameter of 0.2^2 . AWP, which was activated only after the initial 15 epochs, was cited as an absolute necessity for preventing catastrophic overfitting when training over long schedules. The authors noted that removing any single component of this regularization trio led to noticeable performance degradation².

Synthesis: What Transfers to the Curriculum Verifier

For a 100–300 sign known-target verifier operating on mobile devices, the 1D CNN structure transfers directly. A 1.85 million parameter model guarantees a highly efficient inference latency (approximately 17ms on an edge CPU), remaining well within the 200ms real-time constraint². The extraction of only 60–80 highly relevant landmarks (hands, lips, shoulders) combined with zero-filling for occluded coordinates provides a mathematically stable input tensor. The reliance on Swish activations over GELU and causal masking over standard padding directly optimizes the TensorFlow Lite conversion process, preventing the real-time execution graphs from collapsing under dynamic temporal windows².

Section 2: Kaggle: Google – American Sign Language Fingerspelling Recognition (2023)

Fingerspelling operates at a fundamentally different temporal frequency and spatial scale than isolated gloss recognition. The task requires character-level transcription of continuous, highly coarticulated hand micro-movements, where individual letters blur into fluid trajectories. The top solutions universally abandoned fixed-class softmax layers and adopted sequence-to-sequence algorithms typically utilized in Automatic Speech Recognition (ASR).

Per-Source Record: Kaggle ASL Fingerspelling Top Solutions

Group / Author	Year	Task	Dataset	Input Modality	Metrics	Offline vs. Real-Time	Code/Weights Links
Christof Henkel (1st)	2023	Fingerspelling	Kaggle ASL FS	Landmarks	Squeezeformer	Real-Time Capable	³
hoysos48 (2nd Place)	2023	Fingerspelling	Kaggle ASL FS	Landmarks	CTC+Attention	Real-Time Capable	¹⁰

Gezi (3rd Place)	2023	Fingers pelling	Kaggle ASL FS	769-dim features	17-Layer Squeezeformer	Real-Time Capable	¹¹
rob-gilles-رافي (9th)	2023	Fingers pelling	Kaggle ASL FS	Landmarks	10.5M Params	Real-Time Capable	¹²

CTC Variants and Character-Level Modeling

The first-place architecture utilized an improved Squeezeformer encoder paired with a simple 2-layer Transformer decoder³. The second-place solution explicitly cited the adoption of joint Connectionist Temporal Classification (CTC) and Attention models on top of previous ISLR architectures¹⁰.

CTC is mathematically designed to handle unsegmented sequence data where the exact temporal alignment between the input frames and the output characters is unknown. By predicting a probability distribution over the alphabet—including a crucial 'blank' token—at each time step, CTC collapses repeated character predictions into a single output string unless they are separated by the blank token. This directly solves the ASL coarticulation problem where a signer holds a letter for an indeterminate number of frames¹. The reliance on a character-level decoder rather than sub-word tokenization proved essential, as fingerspelling frequently involves arbitrary proper nouns and addresses rather than standard vocabulary.

Inference Efficiency and Temporal Masking

To manage the heavy computational load of deep Squeezeformer networks operating over 384 frames, the third-place architecture introduced a temporal downsampling layer that explicitly reduced the frame count from 320 to 160 early in the network¹¹. This downsampling was necessary to allow the deployment of a deep 17-layer network within the strict compute constraints.

Temporal masking was aggressively utilized across the top solutions. The third-place solution applied a "heavy mask" that entirely zeroed out 50% of the input frames stochastically during training¹¹. This forces the Squeezeformer to learn robust temporal interpolation and rely on the surrounding context of the trajectory, mapping directly to the human visual system's ability to comprehend fingerspelling even when individual frames are blurred. Furthermore, Rotary Position Embedding (ROPE) was utilized over absolute positional encodings because it preserves the relative temporal distance between frames, which is vital for distinguishing the rhythmic cadence of fingerspelling¹¹.

Section 3: PopSign Evidence and Learner Behavior

The PopSign dataset, which served as the foundation for the Kaggle competitions, and the

subsequent analyses surrounding its deployment yield critical insights into how sign language data is captured, curated, and misclassified in real-world educational gaming scenarios¹³.

Per-Source Record: PopSign

Group / Author	Year	Task	Dataset	Input Modality	Metrics	Offline vs. Real-Time	Code/Weights Links
PopSign Team	2023	Data Elicitation	PopSign v1.0	MediaPipe Hands	210,000 videos	Real-Time Game	¹³

Dataset Collection and Preprocessing Constraints

PopSign's foundational dataset comprises over 210,000 video examples captured at a high resolution (1944x2592) from 47 consenting Deaf adult signers whose primary language is ASL¹³. The data was collected entirely via consumer webcams rather than controlled studio environments, establishing a highly realistic baseline for end-user applications. The recognition system relied on MediaPipe Hands to extract 21 key points per hand, applying a strict confidence threshold of 0.5 for detection and 0.1 for tracking¹³.

A load-bearing finding regarding spatial variance is PopSign's handling of handedness. All videos featuring a left-hand dominant signer were mathematically flipped along the vertical axis, forcibly converting them into right-handed coordinate arrays prior to training¹³. Because left-handed signers represent a small minority of the overall population, failing to normalize handedness creates catastrophic class imbalances and spatial confusion within the model.

Error Patterns and Confused Pairs

The manual review of the 210,000 examples revealed that 39,304 examples (nearly 19%) were technically signs but either represented an incorrect lexical variant or were misclassified due to physical similarities¹³.

The dataset documentation highlights specific "confused pairs" that heavily plague coordinate-based recognition systems. The system confused the sign "WAKE" with "AWAKE" in 47.0% of occurrences, and "GLASSWINDOW" with "TOOTH" in 25.1% of occurrences¹³. These confusions underscore the profound limitations of pure kinematic tracking. "WAKE" and "AWAKE" share identical base handshapes and spatial locations, differing only in the micro-tension of the movement parameter and subtle facial grammar. This empirical failure mode demonstrates that known-target validation algorithms cannot rely on a single global confidence threshold; they must explicitly search for the subtle phonological features that separate these minimal pairs to provide accurate verification.

Section 4: Academic Work Closest to Known-Target Validation

An exhaustive survey of the academic literature reveals a profound gap: literature explicitly defining "sign language assessment," "automatic signing evaluation," or "sign language learning feedback" is almost entirely absent. The academic community has focused exclusively on classification (ISLR) and translation (SLT). Therefore, the absence of verification literature dictates that the verification system must be constructed by adapting dictionary retrieval and phoneme-prediction architectures.

Per-Source Record: Academic Validation Context

Group / Author	Year	Task	Dataset	Input Modality	Metrics	Offline vs. Real-Time	Code/Weights Links
Desai et al.	2023	Dictionary Retrieval	ASL Citizen	RGB Video	63% Acc, 91% R@10	Offline	¹⁴
Sem-Lex Team	2023	Phoneme Prediction	Sem-Lex Benchmark	Video / Pose	85% Phoneme Acc	Offline	¹⁶

ASL Citizen and Dictionary Retrieval

The most relevant task framing in the literature is the "dictionary retrieval" paradigm introduced by the ASL Citizen dataset¹⁴. ASL Citizen contains 83,399 videos encompassing 2,731 distinct signs collected from 52 everyday signers¹⁴.

The authors explicitly reframe the traditional ISLR classification task into a retrieval mechanism. When a user demonstrates a sign to a webcam, the algorithm generates an embedding and performs a nearest-neighbor search against a known dictionary. The state-of-the-art baselines published alongside ASL Citizen achieve a top-1 accuracy of 63% and a recall-at-10 metric of 91%¹⁴. Crucially, these metrics are achieved under strict signer-independent evaluation splits, meaning the tested users were entirely unseen during training. The 91% recall-at-10 proves that while extracting the exact intended sign as the top prediction is mathematically difficult across a 2,700-word vocabulary, generating a latent space where the target sign is consistently embedded nearby is highly feasible¹⁴. For a constrained curriculum of 100–300 words, this

retrieval approach serves as a perfect verification backbone.

Sem-Lex and Sub-Lexical Parameter Prediction

The Sem-Lex benchmark provides the architectural foundation for generating per-parameter feedback¹⁶. Sem-Lex aligns 84,000 video productions with the ASL-LEX database, isolating the specific sub-lexical phonemes (handshape, location, movement).

By utilizing a Spatio-Temporal Graph Convolutional Network (SL-GCN), the Sem-Lex researchers demonstrated that individual phonological features can be recognized with 85% accuracy¹⁶. Furthermore, utilizing phoneme prediction as an auxiliary loss target during the training of an ISLR model drastically improves the model's overall spatial representation accuracy. For a known-target verifier, this proves that a properly structured ST-GCN can simultaneously output a global validation score for the sign alongside individual accuracy scores for the learner's handshape and spatial location.

Section 5: Speech CAPT as the Design Template

Because sign language verification literature is sparse, the proposed ASL verification system must inherit its mathematical framework from speech Computer-Assisted Pronunciation Training (CAPT) systems. In spoken language, automated assessment is driven by the Goodness-of-Pronunciation (GOP) algorithm.

Core Scoring Mechanics of GOP

GOP operates by calculating the posterior probability of a target phoneme given the acoustic observation, normalized against the probability of all possible unconstrained phonemes. The system utilizes Hidden Markov Models (HMM) or Deep Neural Networks (DNN) to perform a "forced alignment"—matching the incoming audio frames to the expected sequence of states for the known word.

The standard GOP equation is defined as:

$$GOP(p) = \frac{1}{N} \sum_{t=1}^N \log \frac{P(O_t|p)}{\sum_q P(O_t|q)}$$

Where O_t represents the observation at time frame t , p represents the target phoneme, q represents the set of all possible phonemes, and N is the temporal duration of the phoneme.

Mapping Speech CAPT to ASL Validation

To retrofit GOP for ASL, the acoustic observations (O_t) are replaced by the spatial-temporal landmark matrices (X_t), and the phonetic inventory (p) is replaced by the ASL phonological parameters (ϕ : Handshape, Location, Movement, Palm Orientation).

Speech CAPT Concept	ASL Validation Analog	Operational Definition for App
Acoustic Feature Extraction (MFCCs)	Skeletal Pose Extraction (MediaPipe)	3D coordinate normalization, bone vectors, and velocity derivatives.
Forced Alignment (HMM/Viterbi)	Dynamic Time Warping (DTW) / CTC	Aligning the learner's dynamic temporal frames to the idealized trajectory of the expected target sign.
Phoneme Inventory (Vowels/Consonants)	Sub-Lexical Parameters (ASL-LEX)	The specific handshapes (e.g., '1', 'B', '5') and articulation locations required for the sign.
Target Phone Score	Expected Parameter Probability	Evaluating if the user's handshape matches the expected handshape embedding regardless of overall spatial accuracy.
Unconstrained Phone Loop	Nearest-Neighbor Background Model	Comparing the generated embedding against all 300 signs in the curriculum to ensure the user isn't accidentally performing a different sign.

Feedback Granularity and Failure Modes

Research in speech CAPT explicitly demonstrates that overly granular feedback discourages learners. If a system attempts to correct micro-variations in pitch or tongue placement simultaneously, cognitive overload occurs. For ASL, this implies the validation system should only flag one sub-lexical parameter error at a time (e.g., "Your location is correct, but check your handshape").

The most pervasive failure mode in CAPT is "overstrict scoring" caused by native-speaker bias. If the background model is trained exclusively on highly fluent, rapid native speech, it will ruthlessly reject learners who possess a heavy accent. In ASL, an "accent" manifests as

movement epenthesis, slower velocity, and larger, exaggerated spatial paths. If the known-target verifier's thresholds are calibrated strictly against the 47 fluent Deaf adults from the PopSign dataset, it will incorrectly fail hearing novices even when their semantic execution is perfectly legible.

Section 6: Commercial and Shipped Systems

An audit of deployed, commercial ASL learning applications reveals a stark discrepancy between marketing claims and actual technical capabilities, primarily exposing the immense difficulty of shipping robust verification models to consumer devices.

Per-Source Record: Commercial Systems

Group / Author	Year	Task	Dataset	Input Modality	Metrics	Offline vs. Real-Time	Code/Weight Links
Lingvano	2024	App Delivery	Proprietary	Video/Camera	3 Million Users	Real-Time App	⁵

Lingvano's Marketing vs. Execution

Lingvano is widely regarded as a premier ASL educational app, boasting over 3 million users and high app store ratings¹⁹. The developer's marketing explicitly claims: *"Our sign recognition technology analyzes how you spell each letter and gives you instant feedback. Know instantly if your hand shape and motion are correct!"*⁵.

However, empirical evidence extracted from longitudinal user reviews contradicts this claim for full-word lexical signs. While the app may deploy a localized handshape classifier for static alphabet fingerspelling, multiple user reviews cite the camera feature as functionally inert for conversational validation. As one reviewer explicitly detailed: *"I thought the app would have the camera and make me do that sign and it would recognize what I am doing. It doesn't have that. It does however have the mirror button and it does have the camera, but it doesn't do anything besides show you yourself..."*⁴.

This reveals that Lingvano relies primarily on multiple-choice paradigms, receptive video comprehension, and self-assessment ("mirroring") for expressive signing²⁰.

The UX Failure of Unreliable Recognition

The Lingvano reviews provide a critical UX blueprint for the verifier app. Users actively desire algorithmic feedback but are easily frustrated by rigid pacing¹⁹. Reviewers specifically requested the ability to adjust the speed of the reference signers, noting that Lingvano's native signers execute fingerspelling too rapidly for beginners to parse, despite the existence of a "turtle speed" toggle²⁰.

This confirms the necessity of abstention protocols in the verification system. If a deployed application cannot guarantee a highly calibrated, low-latency validation score, reverting to a "mirror" UX where the user self-evaluates against a reference video is highly preferable to subjecting the learner to a barrage of false negatives generated by an overstrict classifier.

Section 7: Learner vs. Fluent Signing Evidence

To ensure the known-target verification model accommodates novice users, the system must mathematically account for the physical discrepancies between L1 (native) and L2 (learner) signers.

Kinematic Discrepancies

Evidence from dataset collection methodologies, including the ASL Fingerspelling Recognition corpus, specifically targets the capture of "real world speeds" because fluent signing is heavily compressed²².

- **Velocity and Compression:** Fluent signers compress their spatial trajectories, minimizing the physical distance the hand travels between the start and end of a sign to conserve energy. Novices exaggerate the spatial envelope, signing widely outside the standard bounds of the torso.
- **Movement Epenthesis:** Fluent signers seamlessly blend the transitional movements between signs. Novices exhibit distinct hesitation pauses, physically returning their hands to a resting state or stopping abruptly before initiating the next sign.
- **Coarticulation:** In fingerspelling, fluent users merge handshapes (e.g., leaving the thumb extended from the previous letter if it is needed in the subsequent letter). Novices attempt to hit the idealized, textbook posture for every individual character, resulting in a staccato, non-fluid rhythm.

Domain Shift Strategy

To counteract this domain shift, the verification model cannot rely on absolute temporal lengths or raw velocity thresholds to calculate the validation score. The spatial normalization techniques utilized by the Kaggle winners—specifically interpolating the temporal axis to a fixed length (e.g., 160 frames) regardless of the actual execution time—are mathematically mandatory⁶. This process completely isolates the spatial trajectory from the user's execution speed, preventing the system from penalizing a novice purely for signing slowly or pausing.

Synthesis: What Transfers to the Curriculum Verifier

The intersection of Kaggle optimizations, academic retrieval baselines, and speech CAPT theory provides a highly concrete toolkit ready for deployment in a 100–300 sign verifier.

1. **Architecture:** The 1D CNN + Transformer hybrid (under 2 million parameters) from the ISLR Kaggle competition is the optimal spatial embedding extractor. It guarantees sub-200ms latency on edge CPUs and avoids the catastrophic memory overhead of pure vision transformers¹.
2. **Landmark Pruning:** Do not input the entire skeletal graph. Restrict the input tensor to 60–80 highly targeted nodes (predominantly the dominant hand, lips, eyebrows, and

shoulders) to eliminate noise⁶.

3. **Handedness Normalization:** Implement a pre-processing pipeline that detects left-handed signers and vertically flips their coordinates on the X-axis prior to inference, ensuring the model only needs to evaluate right-handed topological maps¹³.
4. **Temporal Masking:** Train the model using 50% temporal masking. This forces the verification system to evaluate the macro-trajectory of the sign rather than over-indexing on individual, static frames that may suffer from webcam motion blur¹¹.
5. **Scoring Paradigm:** Do not use softmax classification. Map the problem to a Goodness-of-Pronunciation metric where the user's generated feature embedding is scored via cosine similarity against the idealized cluster centroid of the known target sign.
6. **Fingerspelling Decoupling:** Utilize a separate Squeezeformer model equipped with a CTC decoder strictly for fingerspelling tasks. CTC is the only proven mathematical method for decoding unsegmented, coarticulated novice spelling³.

Contradictions, Gaps, and Low-Confidence Claims

- **The Verification Literature Gap:** There is zero high-confidence academic literature detailing optimal threshold calibration for ASL validation. The adaptation of speech GOP to ASL topological graphs is a theoretical synthesis drawn from adjacent domains; empirical threshold tuning will require live user testing.
- **3D Depth Ambiguity:** While upstream extractors generate (X, Y, Z) coordinates, the Z-axis (depth) extracted from 2D webcams is notoriously unreliable. Models heavily dependent on distinguishing signs via depth (e.g., pointing toward the camera vs. parallel to the chest) suffer severe accuracy drops.
- **Dataset Bias:** PopSign and ASL Citizen feature predominantly adult signers. Normalizing skeletons based on absolute shoulder width poses an untested risk for children or users with atypical morphologies, potentially leading to immediate verification failure¹.
- **Commercial Capabilities:** Lingvano's claims regarding real-time algorithmic feedback are contradictory. Marketing materials claim instant sign recognition, but user consensus explicitly defines the feature as a passive camera mirror, suggesting that scalable, on-device sign verification remains unsolved in the commercial sector⁴.

Source Table

ID	URL Link	Context
1	uploaded:ASL Pose Representation Research.pdf	Task 2 Specifications and Pose Normalization
2	uploaded:Real-Time ASL Recognition	Task 1 Problem Definition

	Specification.pdf	
6	https://www.kaggle.com/competitions/asl-signs/discussion/406978	Kaggle ISLR 1st Place Code & Inference Speeds
11	https://www.kaggle.com/competitions/asl-signs/discussion?sort=votes	Kaggle ISLR Top Solutions Overview
14	https://www.kaggle.com/competitions/asl-signs/discussion/390935	TFLite Conversion Issues and CNN Optimization
19	https://www.kaggle.com/competitions/asl-signs/writeups/hoyeol-sohn-1st-place-solution-1dcnn-combined-with	Kaggle ISLR 1st Place Architecture Details
20	https://www.kaggle.com/competitions/asl-signs/writeups/hoyeol-sohn-1st-place-solution-1dcnn-combined-with	Kaggle ISLR 1st Place (Secondary Reference)
25	https://www.kaggle.com/competitions/asl-signs/writeups/2nd-place-solution-google-isolated-sign-language-r	Kaggle ISLR 2nd Place Solution Link
27	https://www.kaggle.com/c/asl-fingerspelling/discussion/409438	ASL Fingerspelling Discussion
28	https://www.kaggle.com/c/asl-fingerspelling/discussion/409438	Kaggle ASL Fingerspelling (Secondary Reference)

34	https://www.kaggle.com/competitions/asl-fingerspelling/overview/	Kaggle FS Overview and Dataset Definition
38	https://www.kaggle.com/competitions/asl-fingerspelling/discussion/434588	Kaggle FS 2nd Place ASR and CTC Strategies
43	https://www.kaggle.com/competitions/asl-fingerspelling/writeups/rob-gilles-rafi-9th-place-solution	Kaggle FS 9th Place Conv1D Architecture
46	https://www.kaggle.com/competitions/asl-fingerspelling/writeups/rob-gilles-rafi-9th-place-solution	Kaggle FS 9th Place (Secondary Reference)
60	https://github.com/ChristofHenkel/kaggle-asl-fingerspelling-1st-place-solution	Kaggle FS 1st Place GitHub Repository
63	https://www.kaggle.com/competitions/asl-fingerspelling/writeups/gezi-3rd-place-solution-with-code-17-layers-squeez	Kaggle FS 3rd Place Solution Overview
65	https://www.kaggle.com/competitions/asl-signs/writeups/2nd-place-solution-google-isolated-sign-language-r	Kaggle ISLR 2nd Place Architecture Details
66	https://github.com/ChristofHenkel/kaggle-asl-fingerspelling-1st-place-solution	Kaggle FS 1st Place GitHub (Secondary)
67	https://www.kaggle.com/competitions/asl-signs/writeu	Kaggle ISLR 1st Place

	ps/hoyeol-sohn-1st-place-solution-1dcnn-combined-with	Comprehensive Writeup
70	https://www.kaggle.com/competitions/asl-fingerspelling/writeups/gezi-3rd-place-solution-with-code-17-layers-squeeze	Kaggle FS 3rd Place Temporal Masking Details
74	https://arxiv.org/abs/2304.05934	ASL Citizen Dictionary Retrieval Paper
76	https://www.researchgate.net/publication/369974702_ASL_Citizen_A_Community-Sourced_Dataset_for_Advancing_Isolated_Sign_Language_Recognition	ASL Citizen Paper (Secondary Reference)
77	https://www.microsoft.com/en-us/research/project/asl-citizen/	ASL Citizen Microsoft Project Page
82	https://signdata.cc.gatech.edu/res/doc/popsign_v1_0/popsign_v1_0_supplemental.pdf	PopSign Supplementary Documentation
83	https://signdata.cc.gatech.edu/res/doc/popsign_v1_0/popsign_v1_0_supplemental.pdf	PopSign Data Collection and Minimal Pairs
85	https://par.nsf.gov/servlets/purl/10481698	Sem-Lex Benchmark Paper
90	https://arxiv.org/abs/2310.00196	Sem-Lex Benchmark ArXiv

91	https://apps.apple.com/us/app/lingvano-learn-sign-language/id1547252782?see-all=reviews&platform=ipad	Lingvano App Store Reviews
93	https://apps.apple.com/us/app/lingvano-learn-sign-language/id1547252782?see-all=reviews	Lingvano iOS Reviews on Camera Mirroring
96	https://www.reddit.com/r/asl/comments/1fuuomc/my_thoughts_on_lingvano_after_completing_all_the/	Reddit Learner Reviews on Lingvano
97	https://www.lingvano.com/asl	Lingvano Marketing Copy

Evidence Pack Handoff (For Task 3B)

1. **Architecture Efficiency:** 1D CNNs coupled with lightweight Transformers (under 2 million parameters) are the proven, optimal architectures for isolating signs efficiently on edge devices².
2. **Fingerspelling Decoupling:** Fingerspelling requires a dedicated sequence-to-sequence model utilizing Connectionist Temporal Classification (CTC) to decode fluid, unsegmented micro-movements³.
3. **Temporal Interpolation:** To eliminate the domain shift caused by novices signing much slower than fluent references, temporal landmarks must be interpolated to a fixed sequence length (e.g., 160 frames) prior to classification⁶.
4. **Left-Handed Normalization:** Models will suffer catastrophic class imbalance unless left-handed signers are detected and their spatial coordinates are flipped vertically to match right-handed arrays¹³.
5. **Regularization via Masking:** Randomly masking up to 50% of the temporal frames during training forces the model to learn the macro-trajectory of the sign, preventing overfitting on localized webcam blur¹¹.
6. **Sub-Lexical Accuracy:** Phonological parameters (handshape, location, movement) can be extracted and scored individually with 85% accuracy using Spatio-Temporal Graph Convolutional Networks (ST-GCNs)¹⁶.
7. **The Verification Baseline:** Academic baselines framing ASL recognition as a dictionary retrieval task achieve 91% recall-at-10, proving that generating latent embeddings for

nearest-neighbor verification is highly stable¹⁴.

8. **Minimal Pair Confusions:** Coordinate-based extractors natively fail to distinguish signs with identical base kinematics but different micro-tensions (e.g., WAKE vs. AWAKE), necessitating targeted focal feedback¹³.
9. **Commercial UX Reality:** Despite marketing claims of "instant sign recognition," top-rated commercial apps functionally rely on a passive "mirror" UX for expressive practice, highlighting the immense difficulty of deploying uncalibrated classifiers⁴.
10. **The GOP Scoring Paradigm:** In the complete absence of academic ASL verification literature, the system must mathematically adapt Speech Goodness-of-Pronunciation (GOP), replacing forced phoneme alignment with dynamic time warping of spatial sub-lexical parameters.

Works cited

1. ASL Pose Representation Research.pdf
2. 1st place solution - 1DCNN combined with Transformer - Kaggle, <https://www.kaggle.com/competitions/asl-signs/writeups/hoyeol-sohn-1st-place-solution-1dcnn-combined-with>
3. GitHub - ChristofHenkel/kaggle-asl-fingerspelling-1st-place-solution, <https://github.com/ChristofHenkel/kaggle-asl-fingerspelling-1st-place-solution>
4. Lingvano - Learn Sign Language - Ratings & Reviews - App Store, <https://apps.apple.com/us/app/lingvano-learn-sign-language/id1547252782?see-all=reviews>
5. Lingvano - Learn American Sign Language online, <https://www.lingvano.com/asl>
6. 2nd Place Solution | Google - Isolated Sign Language Recognition, <https://www.kaggle.com/competitions/asl-signs/writeups/2nd-place-solution-google-isolated-sign-language-r>
7. Discussion - Google - Isolated Sign Language Recognition | Kaggle, <https://www.kaggle.com/competitions/asl-signs/discussion?sort=votes>
8. 1st place code with reproducibility - Google - Isolated Sign Language Recognition | Kaggle, <https://www.kaggle.com/competitions/asl-signs/discussion/406978>
9. [Q&A] TFlite conversion (e.g. pytorch to tflite) - Google - Isolated Sign Language Recognition | Kaggle, <https://www.kaggle.com/competitions/asl-signs/discussion/390935>
10. 2nd place solution - Test & Compare ASR algorithms - Kaggle, <https://www.kaggle.com/competitions/asl-fingerspelling/discussion/434588>
11. [3rd place solution with code] 17 layers squeezeformer with timerduce and ROPE - Kaggle, <https://www.kaggle.com/competitions/asl-fingerspelling/writeups/gezi-3rd-place-solution-with-code-17-layers-squeeze>
12. 9th Place Solution | Kaggle, <https://www.kaggle.com/competitions/asl-fingerspelling/writeups/rob-gilles-rafi-9th-place-solution>
13. PopSign ASL v1.0: An Isolated American Sign Language Dataset Collected via

Smartphones,

https://signdata.cc.gatech.edu/res/doc/popsign_v1_0/popsign_v1_0_supplemental.pdf

14. [2304.05934] ASL Citizen: A Community-Sourced Dataset for Advancing Isolated Sign Language Recognition - arXiv, <https://arxiv.org/abs/2304.05934>
15. (PDF) ASL Citizen: A Community-Sourced Dataset for Advancing Isolated Sign Language Recognition - ResearchGate, https://www.researchgate.net/publication/369974702_ASL_Citizen_A_Community-Sourced_Dataset_for_Advancing_Isolated_Sign_Language_Recognition
16. The Sem-Lex Benchmark: Modeling ASL Signs and Their Phonemes - NSF PAR, <https://par.nsf.gov/servlets/purl/10481698>
17. [2310.00196] The Sem-Lex Benchmark: Modeling ASL Signs and Their Phonemes - arXiv, <https://arxiv.org/abs/2310.00196>
18. ASL Citizen - Microsoft Research, <https://www.microsoft.com/en-us/research/project/asl-citizen/>
19. Lingvano - Learn Sign Language - Ratings & Reviews - App Store - Apple, <https://apps.apple.com/us/app/lingvano-learn-sign-language/id1547252782?see-all=reviews&platform=ipad>
20. Lingvano: Sign Language - ASL - Apps on Google Play, <https://play.google.com/store/apps/details?id=com.lingvano.app>
21. My thoughts on Lingvano after completing all the lessons. : r/asl - Reddit, https://www.reddit.com/r/asl/comments/1fuuomc/my_thoughts_on_lingvano_after_completing_all_the/
22. Google - American Sign Language Fingerspelling Recognition | Kaggle, <https://www.kaggle.com/competitions/asl-fingerspelling/overview/>